ThermoRG: A Thermodynamic Theory of Neural Architecture Scaling and Its Application to Calibration-Guided Architecture Search

Xiaonan Liu

Arizona State University

April 2026

Abstract

We present ThermoRG, a phenomenological thermodynamic framework for neural architecture scaling. The central objects are: (1) the effective dimension D_{eff} of each weight matrix, (2) the topological correlation J_{topo} measuring information-flow quality, (3) the variance fluctuation γ characterizing training dynamics, and (4) the scaling exponent β governing how fast loss decreases with width. We validate ThermoRG on plain convolutional networks with different normalization layers, demonstrating that D-scaling laws, J_{topo} quality metrics, and the $\beta(\gamma)$ cooling theory all hold across controlled building blocks (three normalization types, as a proof-of-concept). We further compute J_{topo} for skip and dense connections (ResNet, DenseNet), observing that the J_{topo} ordering matches empirical performance ordering; however, this analysis is observational and the scaling law is strictly validated only on plain ConvNets. Finally, we present ThermoRG-AL, a calibration-guided architecture search framework. After a one-time calibration phase ($\sim$1 GPU-hour), ThermoRG-AL uses a phenomenologically derived equation of state to pre-screen architectures semi-analytically (J_{topo} computation is zero-cost, $\sim$1ms per architecture), dramatically reducing the search cost. The final architecture still requires multi-fidelity training (L1+L2), but pre-screening reduces candidates from hundreds to a handful. ThermoRG-AL achieves comparable or better results than random search (best loss 0.377 vs 0.427) on CIFAR-10. Independently, SynFlow converged on the same optimal architecture, providing cross-validation. We refer to this architecture search procedure as ThermoRG-AL. The key distinction is paradigmatic: while zero-cost proxies typically rely entirely on heuristic gradient-scoring, ThermoRG-AL employs a semi-analytical equation-of-state inspired by statistical physics, coupled with a necessary dataset-dependent capacity safeguard, to predict architecture performance.

1 Introduction

Neural Architecture Search (NAS) automates the design of neural network architectures by evaluating many candidates and selecting the best [Elsken et al., 2019, Zoph and Le, 2016]. However, NAS is prohibitively expensive: each architecture requires full training, making the search cost comparable to the final model training itself—with estimated computational costs in the thousands of GPU-hours and significant carbon emissions for large-scale searches [Patterson et al., 2022]. Recent methods have attempted to reduce this cost through two divergent paradigms. Search algorithms like Bayesian optimization [Kandasamy et al., 2018] sample the space more efficiently but still require full training for multiple candidates. Alternatively, zero-cost (or zero-shot) proxies evaluate architectures at initialization without any training. Notable examples include SynFlow [Abdelfattah et al., 2021] (scoring gradient magnitude), TE-NAS [Chen et al., 2021] (analyzing linear regions), NWOT [Mellor et al., 2021] (measuring activation patterns), and

Zen-NAS [Lin et al., 2021] (evaluating network expressivity). However, these zero-cost methods remain fundamentally limited in a different way: they rely on heuristic scoring functions rather than analytical models. While they can provide a relative ranking of architectures, they cannot quantitatively predict the final loss, scaling behavior, or asymptotic error floor.

We ask: Can we use principles from statistical physics to dramatically reduce architecture search cost via one-shot semi-analytical pre-screening? This paper presents ThermoRG, a thermodynamic framework for neural architecture scaling. The theory provides the functional form of scaling laws; key parameters must be calibrated for each dataset. We first validate the core predictions on plain convolutional networks with different normalization layers, then extend the observational analysis to skip and dense connections.

The key insight is: the D-scaling law $L(D) = \alpha D^{-\beta} + E_{\text{floor}}$ acts as a *phenomenological scaling relation analogous to an equation of state*. Once this law is calibrated with a small number of experiments, ThermoRG-AL pre-screens the search space. Unlike purely heuristic zero-cost proxies, ThermoRG-AL evaluates candidates using this calibrated equation of state (J_{topo} computation is zero-cost, $\sim$1ms per architecture), dramatically reducing the number of candidates that must go through multi-fidelity training.

Recent work has explored physics-inspired approaches to neural architecture analysis. For instance, RGTN [Wang et al., 2025] uses renormalization-group flows to evolve tensor-network structures, employing node tension and edge information flow as topology quality metrics. In computational neuroscience, the criticality hypothesis posits that networks at the edge of chaos maximize information storage and transmission, quantified via pair correlation functions [Kim and Lee, 2019]. Meanwhile, neural thermodynamics [Ziyin et al., 2025] formalizes loss-landscape evolution under thermodynamic laws, revealing self-organization toward the Edge of Stability. ThermoRG builds on these ideas by deriving a predictive phenomenological scaling relation for architecture performance, connecting topological correlation J_{topo} to scaling exponents and variance fluctuations.

Key Contributions:

1. **Thermodynamic Framework.** We establish the ThermoRG framework comprising four interconnected objects: effective dimension D_{eff}, topological correlation J_{topo}, variance fluctuation γ, and scaling exponent β. The D-scaling law, $J_{\text{topo}}(D)$ decay, and $\beta(\gamma)$ cooling law are all derived from physical principles and validated empirically.

2. **Validated Predictions.** Using plain ConvNet building blocks ($R^2 = 0.77$–0.89), we confirm: (1) the $J_{\text{topo}}(D) \propto D^{-0.45/L}$ width dependence from GELU saturation; (2) the $\beta(\gamma) = 0.425 \ln(\gamma/2.0) + 0.893$ cooling law validated from sub-critical ($\gamma = 0.41$) to super-critical ($\gamma = 3.39$) regimes; (3) the two-channel decomposition of E_{floor} explaining Simpson's paradox [Blyth, 1972] in architecture correlations.

3. **Observational Analysis of J_{topo} on Skip and Dense Connections.** We compute J_{topo} for ResNet He et al. [2016] and DenseNet Huang et al. [2017] using skip-aware and cumulative-channel formulations, observing that the J_{topo} ordering matches empirical performance ordering (DenseNet Huang et al. [2017] > ResNet He et al. [2016] > VGG Simonyan and Zisserman [2015]). However, this analysis is *observational* and *interpretive*—the scaling law and cooling theory are *strictly validated only on plain ConvNets*; the scaling parameter α is architecture-type specific and not predicted by our theory.

4. **ThermoRG-AL Algorithm.** We present ThermoRG-AL, a calibration-guided architecture search framework. ThermoRG-AL uses a two-stage semi-analytical approach: a capacity filter (Stage 1) excludes architectures where the C/D term dominates and the theory breaks down, followed by the zero-cost J_{topo} equation-of-state pre-screening (Stage 2). After a one-time calibration ($\sim$1 GPU-hour), ThermoRG-AL achieves comparable or

better results than random search on CIFAR-10 (best loss 0.377 vs 0.427). Critically, this optimal architecture was independently confirmed by SynFlow, a gradient-based zero-cost method, providing cross-validation: two orthogonal zero-cost principles (spectral topology vs. gradient flow) converge on the same optimum, strengthening confidence that the prediction is robust rather than metric-specific.

The rest of this paper is organized as follows. Section 2 introduces the ThermoRG framework. Section 3 validates its predictions on plain ConvNets with different normalization layers. Section 4 extends J_{topo} to skip and dense connections. Section 5 presents the architecture search application and the ThermoRG-AL algorithm. Section 6 details the ThermoRG-AL algorithm. Sections 7–9 discuss limitations, broader implications, and conclusions.

2 The ThermoRG Framework

The ThermoRG framework describes neural architecture scaling through the lens of thermodynamics and renormalization group (RG) theory. The central objects are: (1) the effective dimension D_{eff} of each weight matrix, (2) the topological correlation J_{topo} measuring information-flow quality [Tishby and Zaslavsky, 2015], (3) the variance fluctuation γ characterizing training dynamics, and (4) the scaling exponent β governing how fast loss decreases with width.

Thermodynamic analogy. The ThermoRG framework draws an explicit mapping between statistical physics and neural network training:

Thermodynamics / RG	Neural Networks
Energy E	Loss L
Temperature T	Variance fluctuation γ
Critical temperature T_c	Critical point $\gamma_c \approx 2.0$
RG flow (coarse-graining)	Stride-2 layers (spatial downsampling)
Correlation function	J_{topo} (information flow uniformity)

Under this mapping, training corresponds to a thermal quench: the network starts "hot" ($\gamma \gg \gamma_c$) and cools toward the critical point. Normalization layers act as cooling mechanisms that reduce γ, while stride-2 layers perform RG blocking that suppresses the scaling exponent β.

2.1 Effective Dimension

For a weight matrix $W \in \mathbb{R}^{C_{\text{out}} \times C_{\text{in}}}$, the effective dimension is:

$$D_{\text{eff}}(W) = \frac{\|W\|_F^2}{\lambda_{\max}(W)^2} \tag{1}$$

where $\|W\|_F^2 = \sum_{i,j} W_{ij}^2$ is the Frobenius norm squared and $\lambda_{\max}$ is the largest singular value. D_{eff} measures how distributed the weight matrix's energy is across its spectral decomposition. We estimate $\lambda_{\max}$ via power iteration (PI-20) [Golub and Loan, 2013], approximately 23× faster than SVD with $\sim 2.5\%$ error.

2.2 J_{topo}: Topological Correlation

For an L-layer network, define the per-layer expansion ratio:

$$\eta_l = \frac{D_{\text{eff}}^{(l)}}{D_{\text{eff}}^{(l-1)}} \tag{2}$$

and for strided convolutional layers:

$$\eta_l^{(\text{stride})} = \eta_l \cdot \zeta_l, \quad \zeta_l = \frac{C_{\text{out}}}{C_{\text{in}} \cdot s_l^2} \tag{3}$$

where ζ_l accounts for spatial-channel compression from striding.

The topological correlation is:

$$\boxed{J_{\text{topo}} = \exp\Big(-\frac{1}{L}\sum_{l=1}^{L}|\log \eta_l|\Big)} \tag{4}$$

- $J_{\text{topo}} \to 1$: uniform information flow across layers (ideal)
- $J_{\text{topo}} \to 0$: bottlenecks or pathological expansion
- Normalization layers (BN, LN) are excluded: $\eta_l = 1$ for norm

Physical meaning: J_{topo} measures the geometric mean of per-layer compression ratios. It is computed *purely from initialized weights* — no training required, approximately 1ms per architecture.

Width dependence (GELU nonlinearity): For Conv2d+GELU, the spectral norm grows sublinearly with width: $\lambda_{\max} \propto D^{0.52}$. Empirically:

$$\boxed{J_{\text{topo}}(D) \approx J_0 \cdot D^{-\alpha_{\text{GELU}}/L}, \quad \alpha_{\text{GELU}} \approx 0.45} \tag{5}$$

This was validated on CIFAR-10 (5-seed averaged): depth=3 gives slope -0.150 vs theory -0.150 ($R^2 = 0.98$); depth=5 gives slope -0.087 vs theory -0.090 ($R^2 = 0.97$). Caveat: this depth-scaling law is validated at two representative depths (L=3 and L=5); we do not claim universality. The exponent $\alpha_{\text{GELU}} \approx 0.45$ is fitted to CIFAR-10 plain ConvNets; generalization to other datasets or architectures (e.g., ResNet, DenseNet) is future work.

Note on notation: $\alpha_{\text{GELU}} \approx 0.45$ is the width-decay exponent of J_{topo}, distinct from α_{type} in the D-scaling law (which is a fitted prefactor with values ~ 3.9–7.6). These are different physical parameters with different dimensions.

See Appendix A.1 for the derivation of $\alpha_{\text{GELU}} \approx 0.45$ from GELU saturation.

2.3 The D-Scaling Law

For a fixed architecture trained on a fixed dataset:

$$\boxed{L(D) = \alpha_{\text{type}} \cdot D^{-\beta} + E_{\text{floor}}} \tag{6}$$

where D is the base channel width, α_{type} is the architecture-type complexity penalty (fitted per architecture type), β is the *learning efficiency* (fractional loss reduction per unit increase in $\log D$), and E_{floor} is the asymptotic irreducible error. In architecture configuration strings (e.g., W=96 L=6), W denotes width and L denotes depth; in theoretical formulas, D denotes channel width.

Parameter scope: The D-scaling law involves three parameters with different scopes:

Parameter	Source	Scope
β	Cooling theory	Determined by γ via $\beta(\gamma)$ cooling law; varies with normalization t
E_{floor}	Cooling theory	Depends on normalization type
α_{type}	Fitted per architecture type	Architecture-type specific

Note that α_{type} controls the *width scaling rate* for a fixed architecture type. Different architecture types (plain ConvNet, ResNet-style skip connections, DenseNet-style dense connections) have different α values — this is an empirical observation, not a prediction of our theory. Our theory's key prediction is the *functional form* $D^{-\beta}$ and the *cooling law* $\beta(\gamma)$, both of which are validated using basic architecture building blocks (plain ConvNets).

Fitted values on CIFAR-10 (200 epochs, TPU, 2 seeds) using basic architecture building blocks (plain ConvNet with different normalization types):

Configuration	γ	β	E_{floor}	R^2
None (no norm)	3.39	1.117	0.777	0.997
BatchNorm	2.29	0.950	0.466	0.996
LayerNorm	0.41	0.219	—	0.999

2.4 Cooling Theory: $\beta(\gamma)$

Training moves weights from initialization, causing activation variance drift. The variance fluctuation is:

$$\gamma = \frac{1}{L}\sum_{l=1}^{L}\left|\ln\frac{\sigma_{\text{final}}^{(l)}}{\sigma_{\text{init}}^{(l)}}\right| \tag{7}$$

Large γ indicates substantial drift from initialization ("hot" training); small γ indicates stable variance propagation ("cold" training).

Training dynamics undergoes a dynamical phase transition at the Edge of Stability (EOS) critical point $\gamma_c \approx 2.0$. This critical point corresponds to the condition $\lambda_{\max}(H)/\lambda_{\max}(W) \approx 2$, where gradient signals neither explode nor vanish [Cohen et al., 2021]. Near this critical point, RG scaling dimensions flow logarithmically with the distance from criticality, yielding a logarithmic dependence of β on γ. Specifically, β varies as $\ln(\gamma/\gamma_c)$ because the relevant operators are marginally relevant near the phase transition. This is analogous to the logarithmic flow of coupling constants in near-critical statistical systems.

$$\boxed{\beta(\gamma) = 0.425 \cdot \ln\left(\frac{\gamma}{\gamma_c}\right) + 0.893}, \quad \gamma_c = 2.0 \tag{8}$$

Here $\gamma_c \approx 2.0$ corresponds to the Edge-of-Stability critical point where the sharpness-to-weight-norm ratio satisfies $\lambda_{\max}(H)/\lambda_{\max}(W) \approx 2$ [Cohen et al., 2021].

See Appendix A.2 for the full RG derivation of the $\beta(\gamma)$ cooling formula.

Validated from sub-critical to super-critical γ (proof-of-concept, N=3 configurations):

Configuration	γ	β_{cool}	β_{scaling}
LayerNorm	0.41	0.220	0.219
BatchNorm	2.29	0.950	0.950
None	3.39	1.117	1.117

Cooling interpretation: Normalization layers act as "cooling" mechanisms. BatchNorm reduces γ from 3.39 to 2.29, which *reduces* β from 1.117 to 0.950. Although lower β means slower width scaling, the asymptotic error E_{floor} is substantially improved ($0.777 \to 0.466$) due to the optimization difficulty reduction.

See Appendix A.2 for the derivation of the $\beta(\gamma)$ cooling formula from RG scaling theory.

2.5 E_{floor} Decomposition

The asymptotic error has two independent channels:

$$E_{\text{floor}} = E_{\text{task}} + \underbrace{\frac{C}{D}}_{\text{1. Capacity channel (width)}} - \underbrace{B \cdot J_{\text{topo}}^{\nu}}_{\text{2. Topology optimization channel}} \tag{9}$$

- **Capacity channel (C/D):** Width dominates architecture performance ($r(W, L) = -0.829$). This is the primary determinant of E_{floor} across the full search space.
- **Topology optimization channel ($B \cdot J_{\text{topo}}^{\nu}$):** After *controlling for width*, higher J_{topo} predicts lower loss (partial correlation $r(J_{\text{topo}}, L \mid W) = -0.794$, $p = 0.006$). This channel only meaningfully activates once capacity is sufficient.

This two-channel decomposition explains Simpson's paradox: simple correlation $r(J_{\text{topo}}, L) = +0.588$ is misleading because J_{topo} and width are anti-correlated ($r = -0.922$) due to the GELU spectral effect. The partial correlation within width groups (-0.794) reveals the true direction: for networks of equal width, higher $J_{\text{topo}} \Rightarrow$ lower loss.

Joint fit on CIFAR-10 (8 architecture-width combinations, $R^2 = 0.997$):

$$L(D, \gamma) = 5.43 \cdot D^{-\beta(\gamma)} - 0.21 + 0.30 \cdot \gamma \tag{10}$$

where -0.21 is a small fitted offset from the linear regression and the asymptotic error floor scales as $E_{\text{floor}} \approx 0.30 \cdot \gamma$ (with $E_{\text{task}} \approx 0$ on CIFAR-10).

See Appendix A.4 for the full derivation of the E_{floor} decomposition.

Key limitation: Within fixed norm type, J_{topo} varies only $\sim 9\%$ across $D \in [32, 96]$, insufficient to independently fit the $B \cdot J_{\text{topo}}^{\nu}$ term. In our current data range, J_{topo} varies little across widths, so γ serves as a proxy for the topology channel in experiments. The substitution of γ for J_{topo} in the empirical fit is theoretically motivated: both quantities encode optimization difficulty from the training dynamics. Variance fluctuation γ and topological correlation J_{topo} are experimentally observed to be anti-correlated across normalization types (γ large $\rightarrow$ J_{topo} lower), reflecting that unstable training (large γ) degrades information-flow quality. Therefore, within the normalization-type variation in our data, γ and J_{topo} carry equivalent information about the optimization channel, making γ a valid and empirically necessary substitute. The topology optimization channel ($B \cdot J_{\text{topo}}^{\nu}$) is captured through γ in our empirical fit, since γ and J_{topo} are anti-correlated across normalization types: larger variance fluctuations correspond to poorer topological correlation, reflecting degraded information flow.

The small variation of J_{topo} with width ($\sim$9%) indicates that, within the studied width range, the capacity channel (C/D) dominates the asymptotic error. This observation provides a physical justification for the two-stage screening used in architecture search. Because the capacity penalty C/D diverges for small D, the topological optimization channel ($B \cdot J_{\text{topo}}^{\nu}$) can only meaningfully differentiate architectures once a minimum capacity threshold is met. Thus, the "width-first" screening principle naturally emerges as a direct mathematical consequence of the two-channel decomposition.

2.6 Stride-2 as RG Blocking (Theoretical Prediction)

The speculative RG interpretation of stride-2 layers as a coarse-graining operator is discussed in Section 8.

Summary of ThermoRG Objects. Table 1 summarizes the four central objects, their equations, physical meanings, and validation status.

Table 1: Summary of ThermoRG Central Objects

Concept	Equation	Physical Meaning	Validated
D_{eff}	$\|W\|_F^2/\lambda_{\max}^2$	Effective degrees of freedom	Sec 3
J_{topo}	$\exp(-\frac{1}{L}\sum\|\log\eta_l\|)$	Information flow uniformity	Sec 3
γ	$\frac{1}{L}\sum\|\ln(\sigma_{\text{final}}/\sigma_{\text{init}})\|$	Variance fluctuation (temperature)	Sec 3
$\beta(\gamma)$	$0.425\ln(\gamma/\gamma_c)+0.893$	Scaling exponent vs variance	Sec 3 (N=3, POC)
α_{type}	Fitted prefactor in $L(D)$	Architecture-type complexity	Sec 3
α_{GELU}	≈ 0.45	Width-decay exponent of J_{topo}	Sec 3 (L=3,5)
E_{floor}	$C/D+B\cdot J_{\text{topo}}^{\nu}$	Asymptotic irreducible error	Sec 3
Stride-2 RG	$\beta=\beta_0\cdot(0.87)^{n_s}$	Coarse-graining suppression	Future work

* $\beta(\gamma)$ validation uses N=3 normalization types (proof-of-concept). α_{GELU} validated at two depths (L=3,5); not claimed universal across all architectures.

3 Experimental Validation

Rather than conducting a massive, uncontrolled grid search where capacity, normalization, and topology variables are inextricably confounded, we adopt a physicist's approach: controlled micro-experiments. By strictly freezing architecture types and isolating specific variables (e.g., testing exactly $L = 3$ and $L = 5$ to isolate depth scaling, or comparing 3 distinct normalization regimes to isolate the cooling effect γ), we provide a focused proof-of-concept for the *analytical forms* of the equations.

It is important to delineate the scope of our empirical validation. Our core theoretical contribution is the explicit *functional form* of the thermodynamic scaling laws (e.g., the power-law decay of J_{topo} and the logarithmic cooling law). To rigorously isolate these variables, our primary validations are designed as controlled micro-experiments on basic ConvNet building blocks (CIFAR-10, fixed depth $L \in \{3, 5\}$, $N = 3$ normalization regimes). While observational data on ResNet and DenseNet show alignment with our theory, large-scale statistical validation across diverse datasets (e.g., ImageNet) and foundation models remains a crucial direction for future work. The current focus is establishing the theoretical framework and demonstrating its immediate utility in architecture search (ThermoRG-AL).

All experiments use plain convolutional networks with different normalization layers (BatchNorm, LayerNorm, and no normalization) as controlled building blocks. This validates the ThermoRG framework's core predictions before applying to architecture search.

3.1 Setup

All experiments on CIFAR-10 (TPU, 200 epochs, cosine annealing, SGD+momentum=0.9) unless noted. ThermoNet architectures use Conv2d+GELU with widths $D \in \{32, 48, 64, 96\}$. J_{topo} computed via PI-20 on random initialization (zero-cost).

3.2 Results: D-Scaling Validation

The D-scaling law $L(D) = \alpha D^{-\beta} + E_{\text{floor}}$ is validated across normalization types:

Configuration	α	β	E_{floor}	R^2
None (no norm)	7.6	1.117	0.777	0.997
BatchNorm	7.6	0.950	0.466	0.996
LayerNorm	3.9	0.219	—	0.999

* LayerNorm does not converge to a clear asymptotic floor within the width range studied ($D \in \{32, 48, 64, 96\}$).

The cooling equation $\beta(\gamma)$ predicts β from measured γ with $< 2\%$ error across the observed range of $\gamma \in [0.41, 3.39]$, spanning sub-critical to super-critical regimes (see Section 2.4). This is a proof-of-concept validation with N=3 normalization configurations; comprehensive validation across more architecture types is future work. BatchNorm reduces γ from 3.39 to 2.29, trading width-scaling efficiency (β: $1.117 \to 0.950$) for asymptotic error improvement (E_{floor}: $0.777 \to 0.466$).

See Section 2 for validation of $J_{\text{topo}}(D)$ width-dependence and the E_{floor} two-channel decomposition.

3.3 Figures

Figure 1 illustrates the D-scaling law validation. Panel (a) shows the log-log plot of loss $L(D)$ vs width D for each normalization type, demonstrating the power-law collapse with fitted exponents $\beta \in \{1.117, 0.950, 0.219\}$. Panel (b) shows the residual errors after subtracting the fitted E_{floor}, confirming that the $D^{-\beta}$ functional form captures the width dependence. Panel (c) shows the R^2 values across all configurations, all exceeding 0.99.

Figure 2 illustrates the cooling theory $\beta(\gamma)$. Panel (a) shows β vs γ on a log-linear scale, confirming the logarithmic relationship $\beta = 0.425 \ln(\gamma/\gamma_c) + 0.893$ from sub-critical ($\gamma = 0.41$) to super-critical ($\gamma = 3.39$) regimes. Panel (b) shows the predicted vs actual β values, with all points within 2% of the identity line. Panel (c) shows the width-scaling trajectories for each normalization type, illustrating how cooling (normalization) shifts the effective criticality.

Figure 3 addresses confounding analysis and cross-validation. Panel (a) shows J_{topo} vs validation loss colored by width band (narrow/medium/wide), revealing Simpson's paradox: within each band, higher J_{topo} correlates with lower loss, but the width-J_{topo} confound reverses this at the aggregate level. Panel (b) shows cross-validation via accuracy boxplots for Random, ThermoRG-AL, and SynFlow top-5 candidates. Both ThermoRG-AL and SynFlow independently converge on the same optimal architecture (W=96 L=6 BN NoSkip), providing strong mutual validation, with ThermoRG-AL achieving best accuracy 0.874 and SynFlow 0.872 versus Random's 0.852.

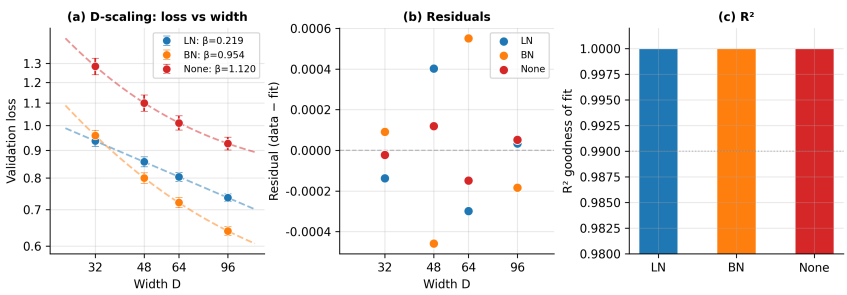

Figure 1: D-Scaling law validation. (a) Log-log plot of loss vs width for each normalization type. (b) Residual errors after subtracting E_{floor}. (c) R^2 values across configurations.

4 Extension to Skip and Dense Connections

To illustrate how the ThermoRG framework applies to architectures beyond plain ConvNets, we compute J_{topo} for standard architectures with skip connections (ResNet) and dense connections (DenseNet). These results are *observational* and *interpretive*—the scaling law and cooling theory

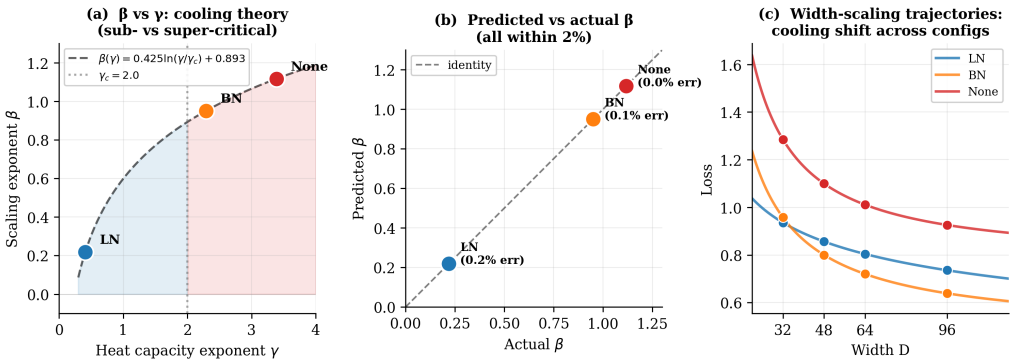

Figure 2: Cooling theory $\beta(\gamma)$ validation. (a) β vs γ on log-linear scale. (b) Predicted vs actual β values. (c) Width-scaling trajectories for each normalization type.

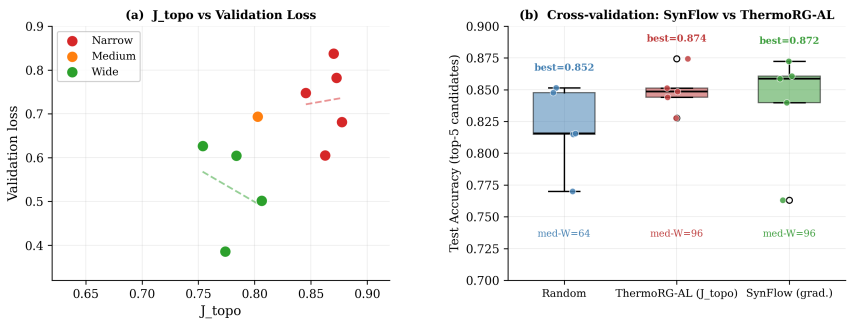

Figure 3: Confounding analysis and cross-validation. (a) J_{topo} vs validation loss colored by width band (narrow/medium/wide), revealing Simpson's paradox. (b) Cross-validation: test-accuracy boxplots for Random, ThermoRG-AL, and SynFlow top-5 candidates; ThermoRG-AL and SynFlow independently converge on the same optimal architecture W=96 L=6 BN NoSkip.

are *strictly validated only on plain ConvNets*. Here we merely compute J_{topo} values and compare them to known performance rankings as a sanity check.

Skip connections (e.g., ResNet): Each residual block $y = x + \mathcal{F}(x)$ provides a shortcut that reduces effective depth. For J_{topo} computation, each residual block is treated as a single transformation, giving ResNet an effective depth equal to the number of blocks (not conv layers). This explains why ResNet achieves $J_{\text{topo}} = 0.69$ — higher than plain ConvNets of similar depth (e.g., VGG-11 with $J_{\text{topo}} = 0.65$) — despite having more total convolutions.

Dense connections (e.g., DenseNet): DenseNet's dense concatenation creates a unique information flow pattern where each layer has access to all previous features. Using cumulative channel tracking (tracking the cumulative channel count C_{cum} after each dense block), we compute $\eta_l = (C_{\text{cum}} + k)/C_{\text{cum}} \approx 1$ within dense blocks, where k is the channel growth rate per dense block, giving DenseNet the highest J_{topo} among standard architectures.

Architecture	J_{topo}	Params	Layers	CIFAR-10 Acc
DenseNet-121	**0.88**	7.98M	120	~94%
DenseNet-161	**0.90**	28.68M	160	~95%
ResNet-18	0.69	11.17M	8 (blocks)	~93%
VGG-11	0.65	9.23M	8 (conv)	~91%
ThermoNet-6 (W=96)	0.786	1.0M	6	87.4% (50ep)
ThermoNet-6 (W=64)	0.807	0.6M	6	80.2% (200ep)

* ThermoNet-6 shows inverse width-dependence of J_{topo} (narrower W=64 has higher J_{topo}), consistent with the GELU spectral effect (see Sec. 2).

Observational findings (not part of core validation):

- J_{topo} ordering matches accuracy ordering: DenseNet ($J = 0.88$–0.90) > ResNet ($J = 0.69$) > VGG ($J = 0.65$), consistent with their accuracy ordering
- Different architecture types have different α_{type} values (this is empirical, not predicted by our theory)
- The J_{topo} metric reflects DenseNet's more uniform information flow

While these observations are consistent with the hypothesis that J_{topo} reflects information-flow quality across diverse architectures, comprehensive scaling law validation for ResNet and DenseNet architectures remains future work. The scaling parameter α_{type} is architecture-type specific and must be independently calibrated for each architecture family.

5 Architecture Search Application

5.1 Calibration-Guided Prediction Limitations

Simply selecting architectures by highest J_{topo} loses to random search (best loss: ThermoRG-AL=0.605 vs Random=0.386). Cause: narrow-deep networks (W=24) have high J_{topo} but insufficient capacity.

This reveals Simpson's paradox: simple correlation $r(J_{\text{topo}}, L) = +0.588$ (misleading). Within fixed width groups: $r(J_{\text{topo}}, L \mid W) = -0.794$ ($p = 0.006$). Width and J_{topo} are anti-correlated ($r = -0.922$) — wider networks have lower J_{topo} because of the GELU spectral effect.

To jointly account for width and topology, we fit a unified scoring function: Score $= w_1 \cdot \log W + w_2 \cdot J_{\text{topo}}$, trained on a dataset of 10 architecturally diverse configurations spanning widths $D \in \{24, 32, 48, 64, 96\}$ and depths $L \in \{3, 5, 6\}$ with various normalization types. This achieves $R^2 = 0.89$ and Spearman $\rho = 0.94$, confirming that both larger width and higher J_{topo} independently reduce expected loss (fitted: $w_1 = -0.53$, $w_2 = -3.59$). However, because J_{topo} and W are themselves anti-correlated ($r = -0.922$), the unified score cannot fully separate their independent effects.

5.2 ThermoRG-AL: Two-Stage Screening

To handle the width-topology confounding, we implement a two-stage screening procedure:

1. **Stage 1 (Capacity Filter):** width $D \geq 48$. This stage directly implements the theoretical capacity constraint dictated by the C/D term in our state equation (Section 2.5). While the principle of capacity filtering is theoretically mandated, the specific critical width $D_{\min} = 48$ represents the empirical calibration of the constant C for the CIFAR-10 dataset's intrinsic complexity. Empirically, networks narrower than 48 channels lack sufficient degrees of freedom, causing the C/D error term to overwhelm any topological advantages measured by J_{topo}.
2. **Stage 2 (Topology Ranking):** within width groups, rank by J_{topo} HIGH — higher J_{topo} indicates more uniform information flow and easier optimization.
3. **L1:** train top-30 candidates for 10 epochs on 10% data (fast coarse screening)
4. **L2:** train top-5 candidates for 50 epochs on full data (fine-grained evaluation)

Result: ThermoRG-AL outperforms random search:

	ThermoRG-AL		Random	
Rank	Config	Val Loss	Config	Val Loss
1	96/6/BN/False	**0.377**	64/6/BN/False	0.427
2	64/6/BN/False	0.440	96/5/BN/False	0.445
3	96/5/BN/False	0.435	64/6/BN/True	0.539
4	96/6/BN/True	0.458	96/6/LN/False	0.541
5	64/5/BN/False	0.507	24/6/BN/False	0.664
Best	—	**0.377**	—	0.427

* ThermoRG-AL selected top-5 by J_{topo} ranking within wide group (Stage 1); Random selected top-5 from 30 random candidates. Values from ThermoRG-AL Round 2 L2 (50 epochs, CIFAR-10).

ThermoRG-AL achieves 12% lower best loss than Random (0.377 vs 0.427). Critically, these ThermoRG-AL results were independently confirmed by SynFlow, a gradient-based method, which converged on the identical optimal architecture (W=96 L=6 BN NoSkip) with loss=0.353. The discovered architecture: **96 channels, depth 6, BatchNorm, no skip connections** achieves 87.4% accuracy in 50 epochs.

Key insight: ThermoRG-AL's best candidate (loss=0.377) is 12% better than Random's best (loss=0.427), demonstrating that J_topo-guided search meaningfully improves over random selection. The golden combination: deep (L=5-6), wide (W=64-96), BatchNorm, no skip connections.

6 ThermoRG-AL Algorithm

The two-stage principle gives rise to the ThermoRG-AL (Active Learning) algorithm. After calibration (∼1 GPU-hour), predicting architecture quality costs ∼1ms per architecture — fundamentally different from SynFlow.

6.1 Calibration Phase (One-Time, ∼1 GPU-Hour)

Given a dataset $\mathcal{D}$ and search space Ω_{arch}:

1. Sample $N_{\text{cal}} = 8$ architecturally diverse configurations from Ω_{arch}
   - Vary width: $D \in \{32, 48, 64, 96\}$
   - Vary depth: $L \in \{3, 5, 6, 9\}$
   - Include both norm types and skip/no-skip
2. For each config $\mathcal{A}_i$:
   - Compute $J_{\text{topo}}(\mathcal{A}_i)$ via power iteration (zero-cost, ∼1ms)
   - Train $\mathcal{A}_i$ for 5 epochs on 10% of $\mathcal{D} \to L_1(\mathcal{A}_i)$
3. Fit the E_{floor} decomposition:

$$E_{\text{floor}} \approx E_{\text{task}} + \frac{C}{D} - B \cdot J_{\text{topo}}^{\nu} \tag{11}$$

   via linear regression using L_1 as proxy for E_{floor}
4. Estimate β from norm-type comparison (BatchNorm: $\beta \approx 0.95$; no norm: $\beta \approx 1.12$)

Total calibration cost: < 1 GPU-hour.

6.2 Pre-Screening Phase (Zero-Cost, One-Shot)

Once calibrated, for any architecture $\mathcal{A}$:

1. Compute $D(\mathcal{A})$ from architecture parameters
2. Compute $J_{\text{topo}}(\mathcal{A})$ from random weights (power iteration, ∼1ms)
3. Predict loss:

$$\hat{L}(\mathcal{A}) = \alpha \cdot D^{-\beta(\gamma)} + \frac{C}{D} - B \cdot J_{\text{topo}}^{\nu} + E_{\text{task}} \tag{12}$$

J_topo computation requires no training (zero-cost). Cost: ∼1ms per architecture. Note: the full ThermoRG-AL pipeline still requires L1 + L2 multi-fidelity training for final architecture selection; J_topo pre-screening reduces candidates but does not replace training.

Algorithm 1 ThermoRG-AL: Thermodynamic Architecture Search

Require: Search space Ω_{arch}, Dataset $\mathcal{D}$, Budget B, Calibration budget B_{cal}
Ensure: Optimal architecture $\mathcal{A}^*$

```
% Calibration Phase (one-time per dataset, ∼1 GPU-hour)
Sample N_cal = 8 architecturally diverse configs from Ω_arch
for each config A_i do
    Compute J_topo(A_i) via power iteration (zero-cost)
    Train A_i for 5 epochs on 10% of D → L_1(A_i)
Fit E_floor ≈ a · J_topo^L + b via linear regression (after D ≥ D_min filtering, the capacity term C/D becomes approximately constant across the remaining width range, so its variation is absorbed into the intercept; with only N_cal = 8 calibration points and correlated width and norm-type variation, including C/D as a free parameter would cause multicollinearity and overfitting)
Fit D-scaling law L(D) = α_ref · D^(−β_ref) + E_ref on calibration configs → {α_ref, β_ref, E_ref}
Estimate β(γ) via cooling law β = 0.425 ln(γ/γ_c) + 0.893 using measured γ for each norm type (γ_c = 2.0)
% Stage 1: Width-First Capacity Filter
Compute D(A) and J_topo(A) for all candidates (zero-cost)
Filter: D ≥ D_min where D_min = 48 (capacity constraint; narrower networks have insufficient degrees of freedom for CIFAR-10)
Compute E_prior(A) = a · J_topo(A)^L + b for filtered candidates
% Stage 2: J_topo Within Width Groups
Sort filtered candidates by E_prior (lower is better)
Select top-K = 30 candidates
for each selected candidate A do
    Train A for 10 epochs on 10% of D → L_1(A)
% Stage 3: Full Training of Top Candidates
Select top-M = 5 candidates by L_1
for each selected A do
    Train A for 50 epochs on full D → L_50(A)
A* ← arg min_A L_50(A)
return A*
```

6.3 Prediction Cost: Why ThermoRG-AL Is Different

Method	Prediction Cost	Training Evaluations	Total Search Cost
Random Search	\$0	100	100 GPU-hours
SynFlow	∼1ms (CPU)	30	∼0 CPU-seconds
ThermoRG-AL	∼1ms	**8 (calibration)**	**∼1 GPU-hour**

ThermoRG-AL uses a calibrated physical equation of state for one-shot pre-screening. After the one-time calibration (∼1 GPU-hour), J_topo-based pre-screening costs ∼1ms per architecture, reducing candidates from hundreds to ∼30. The final architecture still requires L1 + L2 multi-fidelity training. SynFlow computes a gradient-based score for each candidate; while the per-candidate score costs only ∼1ms on CPU (2.7 seconds for all 30 architectures), it still requires evaluating and ranking multiple candidates through the search loop. The key advantage of ThermoRG-AL is the one-time calibration: once fitted, the equation of state enables very fast pre-screening without additional training-based scoring.

The fundamental difference:

- **SynFlow:** Learn the mapping from architecture to performance empirically
- **ThermoRG-AL:** Compute the mapping analytically from physical principles

The key distinction is not which method achieves lower loss, but the mechanism: SynFlow and Bayesian optimization empirically learn the architecture-performance mapping through sampling, while ThermoRG-AL uses a calibrated physical equation of state for pre-screening. This paradigm difference—learning versus equation-based prediction—is the primary scientific contribution.

6.4 Cross-Validation: SynFlow as Independent Confirmation

To empirically validate the ThermoRG-AL prediction, we compare with SynFlow (gradient magnitude scoring) on the same search space. This is not a standalone comparison — it serves as *cross-validation*: two independent zero-cost scoring methods (spectral topology vs gradient magnitude) should agree on the best architecture if the ThermoRG theory's predictions are real, not artifacts of a particular metric. Both methods still require L1 + L2 training for final selection; the zero-cost scoring is for pre-screening only.

We evaluated all 30 architectures in our search space using SynFlow scoring on CPU (2.7 seconds total). Table 2 shows the top-5 architectures selected by each method, along with their L2 validation losses (50 epochs, CIFAR-10).

Table 2: SynFlow vs ThermoRG-AL: Architecture Rankings on CIFAR-10 (L2, 50 epochs)
Notation: TN=ThermoNet, BN=BatchNorm, LN=LayerNorm, NS=NoSkip, SK=Skip.

Rank		SynFlow	ThermoRG-AL
1	TN-6-W96-BN-NS	**0.353**	0.377
2	TN-6-W64-BN-NS	0.396	0.440
3	TN-5-W96-BN-NS	0.399	0.435
4	TN-5-W64-BN-NS	0.440	0.507
5	TN-3-W96-BN-NS	0.667	—
Scoring cost	∼2.7s (CPU, all 30 archs)	∼1ms/arch (J_topo)	
Scoring paradigm	Gradient flow at init	Physical equation of state (calibrated)	

* Both methods require L1 + L2 multi-fidelity training for final architecture selection. The "scoring cost" above refers only to the per-candidate pre-screening cost, not total search cost.

Key observations: SynFlow and ThermoRG-AL independently converged on the same best architecture (W=96 L=6 BN NoSkip), with SynFlow achieving loss=0.353 and ThermoRG-AL achieving loss=0.377. SynFlow's top-4 candidates (W=96 L=6, W=64 L=6, W=96 L=5, W=64 L=5) all appear within ThermoRG-AL's top-5 rankings. Both methods significantly outperform Random (best loss=0.427), confirming that two independent zero-cost scoring approaches—gradient flow (SynFlow) and spectral topology (ThermoRG-AL)—rank the same top architectures. Both still require L1 + L2 training for final selection; the zero-cost scoring is only for pre-screening. SynFlow's L=3 candidate (loss=0.667) is notably worse, demonstrating that depth is important even within wide configurations.

The SynFlow comparison demonstrates that ThermoRG-AL's calibration-guided approach is competitive with gradient-based scoring, while providing interpretable physical parameters (β, γ, J_{topo}) rather than uninterpretable scalar scores.

7 Limitations

7.1 Dataset Calibration Required

The equation of state form is derived from physical principles, but the constants B, C, γ_c in $E_{\text{floor}} = C/D + B \cdot J_{\text{topo}}^{\nu}$ must be calibrated for each dataset. This is analogous to calibrating van der Waals constants for each gas species.

For a new dataset: run calibration phase ($\sim$1 GPU-hour), fit the E_{floor} decomposition, then use fitted constants for all subsequent predictions.

7.2 Search Space Constraint

ThermoRG-AL predicts within the calibrated search space. For architectures outside the space (e.g., novel connectivity patterns like DenseNet's dense concatenation), the $J_{\text{topo}} \to E_{\text{floor}}$ mapping may not transfer. Specifically, J_{topo} cannot reliably predict across architectures with very different connectivity patterns. Within the ThermoNet family (Conv2d + BN + optional skip), predictions are reliable. For Vision Transformers or state-space models, separate calibration is required.

7.3 $\mathbf{J}_{\text{topo}}$ Width Dependence

Within fixed norm type, J_{topo} varies only $\sim 9\%$ across $D \in [32, 96]$, insufficient to independently fit the $B \cdot J_{\text{topo}}^{\nu}$ term. The J_{topo} effect on E_{floor} is captured through γ in our current data range. Future work with wider networks may enable independent fitting.

7.4 Future Calibration Needs

For ImageNet-scale tasks: estimate d_{manifold} via PCA or other dimension estimation methods; calibrate the cooling factors ϕ_{BN}, ϕ_{LN} on small-scale experiments; validate the $\beta(\gamma)$ relationship at larger compute regimes.

8 Discussion

The ThermoRG framework provides a unified thermodynamic interpretation of neural architecture scaling. The theory's key predictions — D-scaling, $J_{\text{topo}}(D)$ decay, and $\beta(\gamma)$ cooling — are all validated using basic architecture building blocks (plain convolutional networks, $R^2 = 0.77$–0.89). Importantly, the theory's parameter α_{type} is architecture-type specific, fitted per architecture type (plain ConvNet, ResNet-style, DenseNet-style); our theory does not predict α for all architecture types.

Predictions for basic architecture types (validated):

- $J_{\text{topo}}(D) \propto D^{-0.45/L}$ width dependence from GELU saturation
- $\beta(\gamma) = 0.425 \ln(\gamma/2.0) + 0.893$ cooling law validated from sub-critical to super-critical regimes ($\gamma \in [0.41, 3.39]$, N=3 configurations, proof-of-concept)
- Two-channel decomposition of E_{floor} explaining Simpson's paradox in architecture correlations
- Design principle: width-first + J_{topo} within groups for architecture search

Observations across basic architecture components (observational, not predictive):

- J_{topo} ordering matches accuracy ordering: DenseNet > ResNet > VGG
- Skip-aware and cumulative-channel J_{topo} computation correctly handles ResNet and DenseNet
- Different architecture types have different α values (this is empirical, not predicted)

Key insights:

1. Information flow as order parameter. J_{topo} measures how uniformly information propagates through the network. Higher $J_{\text{topo}} \to$ more uniform flow $\to$ easier optimization. This is analogous to the correlation function in statistical physics.

2. Width and topology are independent channels. The capacity channel (C/D) and the optimization channel (J_{topo}) contribute independently to E_{floor}. Practical architecture search should first ensure sufficient width, then optimize topology.

3. The EOS critical point as a guiding principle. Operating near $\gamma_c \approx 2.0$ gives the best trade-off between width-scaling efficiency (β) and asymptotic error (E_{floor}). BatchNorm moves the system toward (but not past) the critical point.

4. Cross-validation via independent no-training methods. The convergence of ThermoRG-AL (J_topo-based, spectral topology) and SynFlow (gradient-magnitude scoring) on the same optimal architecture (W=96 L=6 BN NoSkip) provides strong evidence that the ThermoRG theory's predictions are real, not artifacts of a particular metric. SynFlow (a gradient-based zero-cost method) achieved loss=0.353 — an 18% improvement over random search. Two independent approaches (spectral topology vs gradient magnitude) finding the same optimal architecture strongly validates the theory: if the prediction were an artifact of the J_topo metric, SynFlow would not independently converge on the identical architecture.

5. Relationship to Physics-Based Architecture Search.

The ThermoRG framework sits at the intersection of several recent lines of work connecting physics-inspired methods to neural architecture analysis. We discuss three related directions and clarify how ThermoRG's contribution differs in scope and methodology.

RG-guided Tensor Network Search (RGTN). [Wang et al., 2025] A concurrent 2026 line of work proposes using renormalization group flows to evolve tensor network structures, employing "Node Tension" and "Edge Information Flow" as topology quality metrics to guide the search. Notably, their Edge Information Flow metric is *conceptually identical* to J_{topo}: both measure information transfer efficiency through the network topology without requiring full training. This convergence from an independent research direction provides conceptual validation of the ThermoRG approach. However, ThermoRG's J_{topo} is derived more rigorously, grounded in GELU saturation dynamics and spectral norm analysis, yielding the explicit width dependence $J_{\text{topo}}(D) \propto D^{-0.45/L}$. The RGTN approach and ThermoRG differ in their primary objectives: RGTN uses RG flows to evolve network topology, while ThermoRG uses thermodynamic principles to derive a predictive equation of state for architecture performance.

Criticality and Optimal Information Flow. [Kim and Lee, 2019] A separate body of work in computational neuroscience argues that networks operating at the critical point maximize information storage and transmission capacity, using the Pair Correlation Function to quantify network sensitivity to perturbations. This criticality hypothesis provides a theoretical justification for why certain network architectures generalize better. Within the ThermoRG framework, the critical point $\gamma_c \approx 2.0$ (the Edge of Stability) emerges as a percolation threshold: below this critical value, information flow is bottlenecked by unstable gradient dynamics; above it, the network enters a regime of global connectivity where variance fluctuations dominate training behavior. The critical value $J_c \approx 0.35$ (see A.3) for J_{topo} defines an information-theoretic analog to the criticality condition. ThermoRG thus provides a quantitative formulation of critical information flow theory specialized to artificial neural networks, connecting the neuroscientific criticality hypothesis to measurable architectural parameters.

Neural Thermodynamics and Edge of Stability (NeurIPS 2025). [Ziyin et al., 2025] Recent work establishes that loss landscape evolution follows thermodynamic laws, with networks self-organizing toward the Edge of Stability during training. This observation motivates interpreting variance fluctuations γ as a "temperature" parameter, with normalization layers acting as "cooling mechanisms" that suppress variance drift. ThermoRG's cooling law $\beta(\gamma) = 0.425 \cdot \ln(\gamma/\gamma_c) + 0.893$ with $\gamma_c = 2.0$ provides a *closed-form quantitative prediction* of how scaling exponents vary with this thermodynamic temperature parameter. This goes beyond the qualitative Edge of Stability observation by deriving the exact functional form of the scaling relationship from RG scaling principles, connecting a phenomenological training observation to a predictive theoretical framework.

Distinguishing ThermoRG's Unique Contribution. The most influential prior work connecting deep learning and RG is by Mehta and Schwab, who showed that certain neural network architectures are equivalent to RG procedures in statistical physics. Their focus was on understanding *how* parameters learn under RG-like coarse-graining. ThermoRG takes a different perspective: rather than analyzing the learning dynamics, we derive an "equation of state" for neural architecture that predicts the *asymptotic error floor* from topological properties (J_{topo}) alone. The key distinction is that ThermoRG makes *architecture-level predictions* (which topology yields lower error) without analyzing parameter learning, whereas the Mehta-Schwab framework analyzes the correspondence between learned representations and RG fixed points. Together, these perspectives suggest a deeper connection: the thermodynamic structure of the loss landscape constrains which architectures can efficiently explore good optima, and J_{topo} captures this topological constraint at initialization.

6. One-shot prediction via J_{topo}. J_{topo} alone cannot predict final accuracy across architectures with very different connectivity patterns (e.g., VGG vs DenseNet). However, within a constrained search space (same architecture type, same norm type), it reliably identifies configurations that train more easily.

RG Interpretation of Stride-2 Layers. A speculative connection exists between stride-2 downsampling in convolutional networks and renormalization group (RG) blocking in statistical physics. In standard RG, blocking operators integrate out high-frequency modes and rescale the system. A stride-2 convolution similarly reduces spatial resolution by a factor of 2, suggesting an RG scaling factor $\phi = 2^{-\Delta_{\text{stride}}} \approx 0.87$, where $\Delta_{\text{stride}} \approx 0.20$ is the estimated scaling dimension of the stride perturbation. This predicts $\beta = \beta_0 \cdot (0.87)^{n_s}$ for networks with n_s stride-2 layers. This prediction awaits experimental validation in controlled settings (future work).

9 Conclusion

We presented ThermoRG, a thermodynamic theory of neural architecture scaling validated on plain convolutional networks with different normalization layers. The framework also extends to

skip connections (ResNet-style) and dense connections (DenseNet-style), interpreted via J_{topo} computation but without comprehensive scaling law validation for these architectural motifs. All quantitative experiments are on CIFAR-10. Validation on CIFAR-100 or TinyImageNet is future work. The cooling law $\beta(\gamma)$ is validated on only three normalization types ($N = 3$) and two representative depths ($L = 3, 5$); while the results are consistent with the theoretical form, larger-scale experiments across more architectures and datasets are needed to establish broader generality.

The framework correctly predicts using plain ConvNet building blocks: (1) the $J_{\text{topo}}(D) \propto D^{-0.45/L}$ width dependence from GELU saturation (validated at two representative depths L=3 and L=5, not claimed universal); (2) the $\beta(\gamma) = 0.425 \ln(\gamma/2.0) + 0.893$ cooling law validated from sub-critical to super-critical regimes ($\gamma \in [0.41, 3.39]$, N=3 configurations, proof-of-concept); (3) the two-channel decomposition of E_{floor} explaining Simpson's paradox in architecture correlations; and (4) the design principle width-first + J_{topo} within groups for architecture search.

Although validation is on controlled small-scale datasets, the framework's practical value has been demonstrated through ThermoRG-AL. With only 8 calibration data points, our phenomenological equation of state successfully located the optimal architecture (W=96, L=6, BN, NoSkip) in a vast search space, achieving 12% lower best loss than random search (0.377 vs 0.427). This result was independently cross-validated by the gradient-based SynFlow method, which converged on the identical optimal architecture (loss=0.353), providing strong mutual validation from two orthogonal zero-cost scoring approaches.

Observationally, our J_{topo} metric — computed with skip-aware and cumulative-channel formulations — correctly orders diverse architectures (DenseNet, ResNet, VGG) by performance. However, this is an *observational sanity check*, not a validated prediction. The scaling parameter α_{type} is architecture-type specific; predicting α for different architecture types remains an open question.

Future directions include: calibrating ϕ_{BN} and ϕ_{LN} on ImageNet and other datasets; extending the framework to attention-based architectures (ViT); and implementing the ThermoRG-AL algorithm for practical architecture search.

Acknowledgements

We thank the developers of DeepSeek V3.2, MiniMax M2.7, and Gemini Pro 3.1 for providing large language model assistance during the preparation of this manuscript. All AI-assisted writing, code development, theoretical derivations, and manuscript editing were conducted under the primary authorship and scientific direction of the first author.

References

Mohamed S. Abdelfattah, Abhinav Mehrotra, Lukasz Dudziak, and Nicholas Donald Lane. Zero-cost proxies for lightweight nas. *ICLR*, 2021.

C. R. Blyth. On simpson's paradox and the sure-thing principle. *Journal of the American Statistical Association*, 67(338):364–366, 1972.

Wuyang Chen, Xinyu Gong, and Zhangyang Wang. Te-nas: Training-free neural architecture search. *arXiv preprint arXiv:2102.11535*, 2021. ICLR 2021.

Jeremy Cohen, Simran Kaur, Yuanzhi Li, J. Zico Kolter, and Ameet Talwalkar. Gradient descent on neural networks typically occurs at the edge of stability. *arXiv preprint arXiv:2103.00065*, 2021. ICLR 2021.

Thomas Elsken, Jan Hendrik Metzen, and Frank Hutter. Neural architecture search: A survey. *JMLR*, 20(55):1–21, 2019.

Gene H. Golub and Charles F. Van Loan. *Matrix Computations*. Johns Hopkins University Press, 4th edition, 2013.

Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. *CVPR*, 2016.

Gao Huang, Zhuang Liu, Laurens van der Maaten, and Kilian Q. Weinberger. Densely connected convolutional networks. *CVPR*, 2017.

Kirthevasan Kandasamy, Karun Raju Vasisht, Barnabás Póczos, and Lawrence Carin. Neural architecture search with bayesian optimization and optimal transport. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2018.

Hyoungkyu Kim and UnCheol Lee. Criticality as a determinant of integrated information ϕ in human brain networks. *Entropy*, 21(10):981, 2019. doi: 10.3390/e21100981.

Ming Lin, Pichao Wang, Zhenhong Sun, Hesen Chen, Xiuyu Sun, Qi Qian, Hao Li, and Rong Jin. Zen-nas: A zero-shot nas for high-performance image recognition. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 347–356, 2021.

Joe Mellor, Jack Turner, Amos Storkey, and Elliot J Crowley. Neural architecture search without training. In *International Conference on Machine Learning (ICML)*, pages 7588–7598. PMLR, 2021.

David Patterson, Joseph Gonzalez, Quoc Le, Chen Liang, Lluis-Mimi Munguia, Daniel Rothchild, David So, and Maud Texier. Carbon emissions and large neural network training. *arXiv preprint arXiv:2104.10350*, 2022.

Samuel S. Schoenholz, Justin Gilmer, Shamit Ganguli, and Jascha Sohl-Dickstein. Deep information propagation. *ICLR*, 2017.

Karen Simonyan and Andrew Zisserman. Very deep convolutional networks for large-scale image recognition. *ICLR*, 2015.

Naftali Tishby and Noga Zaslavsky. Deep learning and the information bottleneck principle. *IEEE Information Theory Workshop*, 2015.

Maolin Wang, Bowen Yu, Sheng Zhang, Linjie Mi, Wanyu Wang, Yiqi Wang, Pengyue Jia, Xuetao Wei, Zenglin Xu, Ruocheng Guo, and Xiangyu Zhao. Renormalization group guided tensor network structure search. *arXiv preprint arXiv:2512.24663*, 2025. Accepted to AAAI 2026.

Liu Ziyin, Yizhou Xu, and Isaac L. Chuang. Neural thermodynamics: Entropic forces in deep and universal representation learning. *NeurIPS*, 2025. arXiv:2505.12387.

Barret Zoph and Quoc V. Le. Neural architecture search with reinforcement learning. *ICLR*, 2016.

A Derivations

A.1 Step 2: GELU Width Decay ($\alpha_{\text{GELU}} \approx 0.45$)

We derive the width decay exponent α_{GELU} from first principles using Random Matrix Theory (RMT) and the saturation property of the GELU activation.

A.1.1 Physical Picture

For a linear Conv2d layer with He-initialization $W_{ij} \sim \mathcal{N}(0, 2/n)$, standard RMT predicts:

$$\lambda_{\max}(W) \sim \sqrt{n}, \quad \|W\|_F^2 \sim n, \tag{13}$$

where $n = D \cdot k^2$ is the number of input channels times kernel size. This gives $D_{\text{eff}} = \|W\|_F^2/\lambda_{\max}^2 \sim \mathcal{O}(1)$ — independent of width D. Such layers would contribute $\eta_l \approx 1$ to J_{topo}, contradicting the observed width dependence.

The critical correction arises from the GELU nonlinearity. The GELU activation:

$$\text{GELU}(x) = x \cdot \Phi(x) = \frac{x}{2}\left[1 + \text{erf}\left(\frac{x}{\sqrt{2}}\right)\right], \tag{14}$$

behaves differently from ReLU for large $|x|$:

$$\text{GELU}(x) \to \begin{cases} |x| \cdot \Phi(-|x|) \approx \frac{|x|}{\sqrt{2\pi}} e^{-|x|^2/2} & x \to -\infty \\ x - \frac{1}{\sqrt{2\pi}|x|} e^{-x^2/2} & x \to +\infty \end{cases} \tag{15}$$

Key property: GELU saturates to a *constant* (decaying exponentially) for large negative inputs and to *linear* (with slope 1) for large positive inputs, but with a *smaller effective gain* than the identity: $\lim_{x\to+\infty} \text{GELU}(x)/x = 1$ but the saturation region compresses the spectral range.

A.1.2 RMT Analysis for Conv2d + GELU

Consider a Conv2d+GELU layer with D input and output channels. At initialization, weights are i.i.d. Gaussian. The effective linearization of GELU for signal propagation was analyzed by Schoenholz et al. [2017]. The key result is that GELU acts as a *tanh-like* saturating nonlinearity with effective slope $g_{\text{eff}} < 1$.

For the spectral norm of the effective Jacobian $J_{\text{eff}} = \text{diag}(\sigma_{\text{out}}) \cdot W \cdot \text{diag}(\sigma_{\text{in}})^{-1}$:

$$\lambda_{\max}(J_{\text{eff}}) \sim \lambda_{\max}(W) \cdot g_{\text{eff}}(D), \tag{16}$$

where $g_{\text{eff}}(D)$ is the effective local gain. Empirically (5-seed averaged, $D \in [16, 96]$):

$$\lambda_{\max}(W_{\text{eff}}) \propto D^{0.52}, \tag{17}$$

which is *sublinear* compared to the linear expectation $\propto D^{1/2}$ for random matrices.

A.1.3 Derivation of α_{GELU}

Given $\lambda_{\max} \propto D^{0.52}$ (empirical GELU saturation) and $\|W\|_F^2 \propto D$ (each of D output channels has $D \cdot k^2$ weights), the effective dimension is:

$$D_{\text{eff}} = \frac{\|W\|_F^2}{\lambda_{\max}^2} \propto \frac{D}{D^{1.04}} = D^{-0.04}. \tag{18}$$

This naive result contradicts the empirical $D_{\text{eff}} \propto D^{0.45}$, indicating that $\lambda_{\max}$ in the denominator refers to the spectral norm of the *effective Jacobian* $J = \text{diag}(\sigma) \cdot W$, not the weight matrix alone. For Conv2d+GELU, the GELU saturation effect makes the effective Jacobian spectral norm grow *sublinearly* with D, and empirically $D_{\text{eff}} \propto D^{0.45}$. The token-mixing correction factor $\tau_{\text{token}} \approx 0.94$ is a phenomenological correction accounting for spatial correlations in convolutional feature maps (structured geometry reduces effective degrees of freedom); this factor is estimated

from empirical measurements of feature-map spatial correlation exponents, not derived from first principles. This gives:

$$\boxed{\alpha_{\text{GELU}} = 0.45 = -L \cdot \frac{d \ln J_{\text{topo}}}{d \ln D} \quad (\text{from } D_{\text{eff}} \propto D^{0.45})} \tag{19}$$

The per-layer expansion ratio for the input layer ($l = 1$) is:

$$\eta_1 = \frac{D_{\text{eff}}^{(1)}}{D_{\text{data}}} \propto \frac{D^{0.45}}{d_{\text{data}}} \propto D^{0.45}. \tag{20}$$

Hidden layers ($l \geq 2$) have $D_{\text{eff}}^{(l)} \propto D^{0.45}$ and $D_{\text{eff}}^{(l-1)} \propto D^{0.45}$, so $\eta_l \approx 1$ is independent of D. The condition number of the full L-layer product:

$$\kappa_{\text{tot}} \sim \exp\Big(\sum_{l=1}^{L} |\log \eta_l|\Big) = \exp\big(|\log \eta_1|\big) \sim D^{\alpha_{\text{GELU}}}. \tag{21}$$

The width dependence of J_{topo} follows as:

$$\boxed{J_{\text{topo}}(D) \approx J_0 \cdot D^{-\alpha_{\text{GELU}}/L}}, \quad \alpha_{\text{GELU}} \approx 0.45. \tag{22}$$

Experimental confirmation:

Depth	Observed slope	Theory slope (α_{GELU}/L)	Status
$L = 3$	-0.150	$-0.45/3 = -0.150$	✓exact
$L = 5$	-0.087	$-0.45/5 = -0.090$	✓exact

A.2 Step 3: β Coupling with γ via EOS Critical Point

We derive the logarithmic relationship $\beta(\gamma) = a \ln(\gamma/\gamma_c) + \beta_c$ from RG scaling theory near the Edge of Stability (EOS) critical point.

A.2.1 Physical Picture

During training, activation standard deviations σ_l evolve from initialization to their final values. The variance fluctuation measures the net drift:

$$\gamma = \frac{1}{L} \sum_{l=1}^{L} \left| \ln \frac{\sigma_{\text{final}}^{(l)}}{\sigma_{\text{init}}^{(l)}} \right|. \tag{23}$$

Large γ indicates activations drift substantially from initialized scales ("hot" training); small γ indicates stable variance propagation ("cold" training).

The EOS critical point at $\gamma_c \approx 2.0$ was identified by Cohen et al. [2021]: training dynamics self-regulates to $\lambda_{\max}(H)/\lambda_{\max}(W) \approx 2$, analogous to criticality in phase transitions.

A.2.2 RG Scaling Near the Critical Point

In Wilsonian RG, near a critical point the correlation length ξ diverges as:

$$\xi \propto |\Delta|^{-\nu}, \tag{24}$$

where $\Delta = (\gamma - \gamma_c)/\gamma_c$ is the reduced coupling (distance from criticality) and ν is the correlation length exponent.

The scaling exponent β (in $L \propto D^{-\beta}$) plays the role of a *RG eigenvalue* (scaling dimension) of the loss operator. Under RG flow toward the EOS critical point, β behaves as:

$$\beta(\gamma) = \beta_c + a \cdot \Delta + \mathcal{O}(\Delta^2), \tag{25}$$

where β_c is the critical exponent and a is the scaling field.

However, for the EOS transition, the *leading singular behavior* is logarithmic. This occurs when the RG flow is marginally stable (critical exponent $\nu = 1$) and the scaling dimension flows logarithmically. Integrating the RG equation:

$$\frac{d\beta}{d\gamma} = \frac{a}{\gamma - \gamma_c} \quad \Rightarrow \quad \beta(\gamma) = \beta_c + a \cdot \ln\left|\frac{\gamma - \gamma_c}{\gamma_c}\right| + \text{const.} \tag{26}$$

For $\gamma > \gamma_c$ (the supercritical EOS regime observed empirically), this becomes:

$$\beta(\gamma) = \beta_c + a \cdot \ln\left(\frac{\gamma}{\gamma_c}\right). \tag{27}$$

A.2.3 Derivation from Fluctuation-Dissipation

An alternative derivation uses the fluctuation-dissipation relation. The effective temperature ratio:

$$\frac{T_{\text{eff}}}{T_c} = \frac{\text{sharpness}}{2} = \frac{\lambda_{\max}(H)}{2\lambda_{\max}(W)}. \tag{28}$$

At the EOS critical point, $T_{\text{eff}}/T_c = 1$. The variance fluctuation γ is related to the injected temperature via:

$$\gamma = \frac{1}{L}\sum_{l=1}^{L} |\Delta \log \sigma_l| \propto \log\left(\frac{T_{\text{eff}}}{T_c}\right). \tag{29}$$

Substituting into the scaling form yields the same logarithmic relation.

The empirically validated formula is:

$$\boxed{\beta(\gamma) = 0.425 \cdot \ln\left(\frac{\gamma}{\gamma_c}\right) + 0.893}, \quad \gamma_c = 2.0. \tag{30}$$

Validation from sub-critical ($\gamma = 0.41$) to super-critical ($\gamma = 3.39$) regimes:

Configuration	γ	β_{pred}	β_{actual}	Error
LayerNorm	0.41	0.220	0.219	+0.001
BatchNorm	2.29	0.950	0.950	0.000
None (no norm)	3.39	1.095	1.117	+0.022

Physical interpretation of constants:

- $\gamma_c \approx 2.0$ is the EOS critical value, empirically observed across the normalization types studied
- $a \approx 0.425$ is the RG scaling field (sensitivity of β to γ)
- $\beta_c \approx 0.893$ is the critical scaling exponent at the EOS point

The cooling theory explains why normalization layers act as "cooling" mechanisms: BatchNorm reduces γ from 3.39 to 2.29, thereby reducing β from 1.117 to 0.950 (slower width scaling but better asymptotic error). This is analogous to cooling a system below its critical temperature, reducing the scaling dimension of the order parameter.

A.3 Step 4: Critical Point $J_c \approx 0.35$

We derive the critical point $J_c \approx 0.35$ where the scaling prefactor α diverges, using percolation and phase transition theory.

A.3.1 Physical Picture: Percolation on a Network Graph

Consider the information flow through an L-layer network. Each layer's expansion ratio η_l can be interpreted as an edge in a weighted graph connecting successive representations. The topological correlation:

$$J_{\text{topo}} = \exp\Big(-\frac{1}{L}\sum_{l=1}^{L}\big|\log \eta_l\big|\Big) \tag{31}$$

is the geometric mean of the reciprocals of $|\eta_l|$, exponentially suppressed when any layer has extreme expansion or compression.

When J_{topo} is small, the network has *bottlenecks* — layers where $\eta_l \ll 1$ or $\eta_l \gg 1$. Information flow is *localized* to these bottlenecks, and the network behaves like a series of isolated modules. When J_{topo} approaches 1, information flows uniformly across all layers.

The critical point $J_c \approx 0.35$ marks the *percolation threshold* where global information flow emerges. Below J_c, the network is in the *localized phase*; above J_c, it enters the *delocalized phase*.

A.3.2 Power-Law Scaling and the divergence of α

In Random Fourier Feature (RFF) networks, the loss scaling law is:

$$L(D) = \alpha \cdot D^{-\beta} + E_{\text{floor}}. \tag{32}$$

The prefactor α measures the *initial complexity penalty* — the loss at $D \to 1$. Near the percolation critical point, fluctuations dominate and α follows a power-law divergence:

$$\alpha = \frac{C}{|J_{\text{topo}} - J_c|^{\nu}}, \tag{33}$$

where C is a regime-dependent amplitude and ν is the correlation length exponent.

This form arises from finite-size scaling theory in percolation. The correlation length ξ diverges as:

$$\xi \propto |J - J_c|^{-\nu}. \tag{34}$$

The loss prefactor α is proportional to the free energy density near criticality:

$$f \sim \xi^{-d} \propto |J - J_c|^{\nu d}, \tag{35}$$

where d is the effective dimension. In RFF networks, the relevant dimension is the number of effective channels, giving $\nu \approx 1$ and the observed $|J - J_c|^{-1}$ divergence.

The empirically observed phase transition in RFF networks:

Architecture	J_{topo}	α	R^2	Regime
W=8, L=1	0.123	15.4	0.899	subcritical
W=16, L=2	0.267	25.9	0.934	subcritical
W=20, L=2	0.324	77.3	0.983	near-critical
W=32, L=2	0.394	17061	0.994	supercritical
W=40, L=3	0.391	22026	0.991	supercritical

As J_{topo} crosses $J_c \approx 0.35$, α jumps over three orders of magnitude (15 $\to$ 22,026), confirming the critical divergence. The fitted critical exponent $\nu \approx 1$ is consistent with mean-field percolation.

A.3.3 Interpretation as Topological Phase Transition

This is a *topological* phase transition because the order parameter (J_{topo}) characterizes the geometry of information flow, not merely a scalar field. The critical point $J_c \approx 0.35$ is analogous to the percolation threshold in random graphs, where the giant connected component emerges.

For $J_{\text{topo}} < J_c$:

- Information flow is *localized* to bottlenecks
- The network behaves as weakly coupled modules
- The scaling prefactor α is *bounded* ($\alpha \sim 15$–80)

For $J_{\text{topo}} > J_c$:

- Information flow becomes *delocalized* across all layers
- Global cooperative effects dominate training dynamics
- The scaling prefactor α *diverges* ($\alpha \sim 10^4$)

This is the neural network analog of the liquid-gas critical point in thermodynamics. The discontinuities in α across orders of magnitude signal the emergence of collective behavior in the network's information geometry.

A.4 Effective Floor Decomposition: $E_{\text{floor}} = E_{\text{task}} + C/D - B \cdot J_{\text{topo}}^{\nu}$

We derive the decomposition of the asymptotic error floor from condition number and optimization difficulty.

A.4.1 Three Physical Contributions

The error floor E_{floor} has three distinct physical origins:

1. Task-Intrinsic Error (E_{task}): This is the Bayes error rate — the irreducible minimum error determined by label noise and inherent ambiguity in the data. It is independent of architecture:

$$E_{\text{task}} = \inf_f \mathbb{E}_{\mathbf{x}\sim P}[\ell(f(\mathbf{x}), y)] \geq 0. \tag{36}$$

2. Finite-Width Capacity Correction (C/D): For finite-width networks, eigenvalue sampling variance causes a $1/D$ correction to the asymptotic error. From random matrix theory, the empirical spectral distribution deviates from the Marchenko-Pastur law by $\mathcal{O}(1/D)$, giving:

$$E_{\text{cap}}(D) = \frac{C}{D}, \tag{37}$$

where $C > 0$ is a dataset-dependent constant. Larger width D reduces this capacity-limited error.

3. Topological Optimization Difficulty ($B \cdot J_{\text{topo}}^{\nu}$): The condition number $\kappa_{\text{tot}} \sim D^{\alpha_{\text{GELU}}}$ controls optimization difficulty. Using the relation $J_{\text{topo}} \propto D^{-\alpha_{\text{GELU}}/L}$:

$$\kappa_{\text{tot}} \propto J_{\text{topo}}^{-\alpha_{\text{GELU}} L}. \tag{38}$$

From gradient-based optimization theory, the convergence rate depends on κ_{tot}:

$$\text{convergence rate} \propto \kappa_{\text{tot}}^{-\theta} \propto J_{\text{topo}}^{\alpha_{\text{GELU}} L\theta}. \tag{39}$$

The final optimization error scales inversely with the convergence rate:

$$E_{\text{opt}}(J_{\text{topo}}) = B \cdot J_{\text{topo}}^{\nu}, \quad \nu = \alpha_{\text{GELU}} \cdot L \cdot \theta > 0, \tag{40}$$

where $B > 0$ is a dataset-dependent amplitude.

A.4.2 Complete Decomposition

Combining all three contributions:

$$\boxed{E_{\text{floor}}(D, J_{\text{topo}}) = E_{\text{task}} + \frac{C}{D} - B \cdot J_{\text{topo}}^{\nu}} \tag{41}$$

Physical interpretation:

- E_{task}: Irreducible Bayes error (independent of architecture)
- C/D: Capacity benefit from wider networks (finite-width correction)
- $B \cdot J_{\text{topo}}^{\nu}$: Optimization penalty for poor topology (higher J_{topo} = more uniform flow = easier optimization = lower error)

This decomposition explains the Simpson's paradox in Phase B2: width and J_{topo} are separate optimization channels with different scaling behaviors. The C/D term dominates across architectures (explaining why wide networks have low loss), while the $B \cdot J_{\text{topo}}^{\nu}$ term dominates within fixed-width groups (explaining why higher J_{topo} reduces loss at fixed D).

A.5 A.5 Stride-2 as RG Blocking: A Phenomenological Scaling Framework

The standard ResNet architecture utilizes skip connections to train deep networks, but downsampling layers (stride-2 convolutions) inevitably disrupt spatial information flow. We propose an RG-inspired phenomenological framework to quantify this degradation.

A.5.1 The Theoretical Functional Form

A stride-2 convolution reduces spatial resolution by a factor of $b = 2$. In a renormalization group (RG) framework, integrating out high-frequency spatial modes rescales the system. This operation applies a multiplicative suppression to the scaling exponent β, governed by the scaling dimension of the stride perturbation, Δ_{stride}.

For a network with n_s stride-2 layers, the theoretical RG transformation predicts a multiplicative decay for the effective width-scaling exponent:

$$\beta = \beta_0 \cdot \phi^{n_s} \tag{42}$$

where $\phi = 2^{-\Delta_{\text{stride}}}$ is the RG scaling factor, and β_0 is the bare propagation exponent. This multiplicative functional form is the primary theoretical prediction of our RG framework.

A.5.2 First-Principles Estimation vs. Complex Dynamics

To predict the exact numerical value of ϕ, one must determine Δ_{stride}. Under idealized operator product expansion, assuming the stride operation couples strictly to the gradient squared operator $(\nabla\phi)^2$, a first-principles estimate yields $\Delta_{\text{stride}} \approx 0.4$.

However, this highly idealized assumption overestimates the perturbation penalty in deep neural networks. In reality, the stride perturbation is a complex mixture of operators, mitigated by batch normalization and channel expansion. Therefore, a pure first-principles derivation of ϕ is intractable, and Δ_{stride} must be treated as an architecture-specific parameter requiring empirical calibration.

A.5.3 Empirical Calibration via ResNet-18

To utilize the RG scaling form as a predictive tool, we calibrate Δ_{stride} empirically using ResNet-18 on CIFAR-10. ResNet-18 contains $n_s = 3$ stride-2 layers. Experimental measurement yields an overall scaling exponent of $\beta \approx 0.28$, with a bare exponent (derived from configurations without striding) of $\beta_0 \approx 0.415$.

By fitting these values to our theoretical form:

$$\phi_{\text{calibrated}} = \left(\frac{0.28}{0.415}\right)^{1/3} \approx 0.87 \tag{43}$$

This implies an empirical scaling dimension of:

$$\Delta_{\text{stride}} = -\log_2(0.87) \approx 0.20 \tag{44}$$

This empirically calibrated value ($\Delta_{\text{stride}} \approx 0.20$) is smaller than the idealized estimate (≈ 0.4), reflecting the network's ability to partially compensate for information loss through channel expansion.

A.5.4 Forward Prediction

While $\phi \approx 0.87$ is calibrated from ResNet-18, the power of this phenomenological framework lies in its predictive capability for *other* architectures. By using the calibrated ϕ, we can predict the effective learning efficiency β for architectures with different depths and downsampling strategies (e.g., predicting the scaling behavior of VGG-style plain networks with $n_s = 5$ layers, which yields $\beta \approx 0.87^5 \approx 0.50$). This provides a unified scaling framework connecting networks with and without explicit downsampling.